

Ruvims Ungeitis, st84951, 4301bna,Variant 2

Introduction

1. Import tools and file

```
In [1]: # installing necessary libraries
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from scipy import stats
import statsmodels.formula.api as smf
from statsmodels.stats.outliers_influence import OLSInfluence
from statsmodels.stats.diagnostic import het_breuschpagan, het_goldfeldqu
from statsmodels.stats.stattools import durbin_watson

# loading the data
data = pd.read_excel('ChangedFile.xlsx')
data.head()
```

```
Out[1]:
```

	Rate	Survived	Total	Failed	new_column	new_column2
0	1	yes	7287	1	166.260	166.260
1	2	no	4634	0	136.987	136.987
2	3	yes	5769	1	184.101	184.101
3	3	yes	1870	1	130.129	130.129
4	2	no	3697	0	156.878	156.878

1.2 Data check

```
In [2]: data.info()
data.describe()
data.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Rate            100 non-null    int64
1   Survived        100 non-null    object
2   Total           100 non-null    int64
3   Failed          100 non-null    int64
4   new_column      100 non-null    float64
5   new_column2     93 non-null     float64
dtypes: float64(2), int64(3), object(1)
memory usage: 4.8+ KB
```

```
Out[2]: Rate            0
        Survived       0
        Total          0
        Failed         0
        new_column     0
        new_column2    7
        dtype: int64
```

1.3 Dropping nulls

```
In [3]: newTable = data.dropna()
        newTable.info()
        newTable.describe()
        newTable.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 93 entries, 0 to 99
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Rate            93 non-null    int64
1   Survived        93 non-null    object
2   Total           93 non-null    int64
3   Failed          93 non-null    int64
4   new_column      93 non-null    float64
5   new_column2     93 non-null    float64
dtypes: float64(2), int64(3), object(1)
memory usage: 5.1+ KB
```

```
Out[3]: Rate            0
        Survived       0
        Total          0
        Failed         0
        new_column     0
        new_column2    0
        dtype: int64
```

2. Construct the multipme regression and analyse

2.1 Models quality parameters and t-test

```
In [4]: y = newTable['Rate']
x = newTable[['Total', 'new_column', 'new_column2']]

model = smf.ols("Rate ~ Total + new_column + new_column2", data=newTable)

print(model.summary())
```

OLS Regression Results								
=====								
=====								
Dep. Variable:		Rate		R-squared:				
0.057								
Model:		OLS		Adj. R-squared:				
0.036								
Method:		Least Squares		F-statistic:				
2.742								
Date:		Wed, 20 May 2026		Prob (F-statistic):		0.		
0698								
Time:		14:45:13		Log-Likelihood:		-16		
4.81								
No. Observations:		93		AIC:		3		
35.6								
Df Residuals:		90		BIC:		3		
43.2								
Df Model:		2						
Covariance Type:		nonrobust						
=====								
=====								
	coef	std err	t	P> t	[0.025			
0.975]								

Intercept	2.9101	0.709	4.106	0.000	1.502			
4.318								
Total	-0.0001	5.39e-05	-2.087	0.040	-0.000	-5.4		
2e-06								
new_column	0.0025	0.002	1.099	0.275	-0.002			
0.007								
new_column2	0.0025	0.002	1.099	0.275	-0.002			
0.007								
=====								
=====								
Omnibus:		1.725		Durbin-Watson:				
2.067								
Prob(Omnibus):		0.422		Jarque-Bera (JB):				
1.733								
Skew:		0.314		Prob(JB):				
0.420								
Kurtosis:		2.769		Cond. No.		8.46		
e+17								
=====								
=====								
=====								

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The smallest eigenvalue is 4.58e-27. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

The initial model has low explanatory power. Total is statistically significant, while new_column and new_column2 are not. The model also shows possible multicollinearity, probably because new_column and new_column2 contain very similar information.

2.2 Residual analysis (Gauss-Markov condition)

```
In [5]: residuals = model.resid
fitted_values = model.fittedvalues


plt.figure(figsize=(10, 6))
plt.scatter(fitted_values, residuals, alpha=0.5)
plt.axhline(0, color='red', linestyle='--')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.title('Residuals vs Fitted Values')
plt.show()
```



The residuals are located around the zero line, but the plot shows a slight pattern and the spread becomes larger for some fitted values. This may indicate non-constant variance. There is also at least one extreme point, so outliers should be checked.

Check normalisation by Q-Q plot

```
In [6]: plt.figure(figsize=(6, 6))
stats.probplot(residuals, dist="norm", plot=plt)
plt.title('Q-Q Plot of Residuals')
plt.show()
```



The residuals are approximately normal because most points follow the diagonal line. However, there is some deviation in the tail, which may indicate an outlier.

3. Find outliers


```
In [7]: # Calculate standardized residuals
influence = OLSInfluence(model)
standardized_residuals = influence.resid_studentized_internal

# Calculate Cook's distance
cooks_d = influence.cooks_distance[0]

# Plot Cook's distance
plt.figure(figsize=(10, 6))
plt.stem(np.arange(len(cooks_d)), cooks_d, markerfmt=","")
plt.axhline(4 / len(cooks_d), color='red', linestyle='--', label="Threshold")
plt.xlabel('Observation Index')
plt.ylabel('Cook's Distance')
plt.title("Cook's Distance for Outlier Detection")
plt.legend()
plt.show()

# Identify potential outliers
```

```
outliers = np.where((np.abs(standardized_residuals) > 2) | (cooks_d > 4 /
print(f"Potential outliers at indices: {outliers}")
```



Potential outliers at indices: [6 9 36 51 53 75 78]

4.1 removing outliers

```
In [8]: clean_table = newTable.copy()
removed_outliers = []

while True:
    current_model = smf.ols(
        "Rate ~ Total + new_column + new_column2",
        data=clean_table
    ).fit()

    influence = OLSInfluence(current_model)
    standardized_residuals = np.asarray(influence.resid_studentized_inter
    cooks_d_current = np.asarray(influence.cooks_distance[0])

    threshold = 4 / len(clean_table)

    outliers_current = np.where(
        (np.abs(standardized_residuals) > 2) |
        (cooks_d_current > threshold)
    )[0]

    if len(outliers_current) == 0:
        break

    largest_outlier_position = outliers_current[np.argmax(cooks_d_current
    largest_outlier_index = clean_table.index[largest_outlier_position]

    removed_outliers.append({
        "position": int(largest_outlier_position),
        "index": int(largest_outlier_index),
```

```

        "cooks_distance": float(cooks_d_current[largest_outlier_position])
        "standardized_residual": float(standardized_residuals[largest_out
    })

    clean_table = clean_table.drop(index=largest_outlier_index)

print("Removed outliers:")
for i, item in enumerate(removed_outliers, start=1):
    print(
        f"{i}. position={item['position']}, "
        f"index={item['index']}, "
        f"Cook's distance={item['cooks_distance']:.4f}, "
        f"standardized residual={item['standardized_residual']:.4f}"
    )

print(f"Rows before removing outliers: {len(newTable)}")
print(f"Rows after removing outliers: {len(clean_table)}")

plt.figure(figsize=(10, 6))
plt.stem(np.arange(len(cooks_d_current)), cooks_d_current, markerfmt=","")
plt.axhline(4 / len(cooks_d_current), color='red', linestyle='--', label=
plt.xlabel('Observation Index')
plt.ylabel("Cook's Distance")
plt.title("Cook's Distance for Outlier Detection")
plt.legend()
plt.show()


```

Removed outliers:

1. position=78, index=84, Cook's distance=0.1307, standardized residual=2.4334
2. position=53, index=59, Cook's distance=0.0690, standardized residual=1.7515
3. position=74, index=81, Cook's distance=0.0412, standardized residual=2.8488
4. position=51, index=57, Cook's distance=0.0446, standardized residual=-1.2237
5. position=6, index=7, Cook's distance=0.0514, standardized residual=-1.2858
6. position=14, index=16, Cook's distance=0.0475, standardized residual=2.0310
7. position=29, index=35, Cook's distance=0.0376, standardized residual=2.0564
8. position=31, index=38, Cook's distance=0.0291, standardized residual=2.1212
9. position=67, index=79, Cook's distance=0.0294, standardized residual=2.0682
10. position=8, index=10, Cook's distance=0.0206, standardized residual=2.5239
11. position=20, index=25, Cook's distance=0.0337, standardized residual=-2.0426
12. position=30, index=40, Cook's distance=0.0221, standardized residual=-2.5193
13. position=75, index=94, Cook's distance=0.0555, standardized residual=-1.5363

Rows before removing outliers: 93

Rows after removing outliers: 80



Outliers were removed one by one. At each step, the model was rebuilt, Cook's distance was recalculated, and only the most influential remaining observation was removed.

```
In [9]: model_clean = smf.ols(  
        "Rate ~ Total + new_column + new_column2",  
        data=clean_table  
    ).fit()  
  
    print(model_clean.summary())
```



```

=====
                        OLS Regression Results
=====
=====
Dep. Variable:          Rate      R-squared:
0.115
Model:                  OLS      Adj. R-squared:
0.092
Method:                 Least Squares      F-statistic:
4.982
Date:                   Wed, 20 May 2026      Prob (F-statistic):          0.0
0923
Time:                   14:45:13      Log-Likelihood:          -11
9.90
No. Observations:      80      AIC:          2
45.8
Df Residuals:          77      BIC:          2
52.9
Df Model:              2
Covariance Type:       nonrobust
=====
=====
                        coef      std err          t      P>|t|      [0.025
0.975]
-----
Intercept              2.3899      0.622        3.843      0.000        1.152
3.628
Total                 -0.0001      4.57e-05      -2.492      0.015      -0.000   -2.2
9e-05
new_column             0.0039      0.002        1.909      0.060      -0.000
0.008
new_column2            0.0039      0.002        1.909      0.060      -0.000
0.008
=====
=====
Omnibus:              6.504      Durbin-Watson:
1.802
Prob(Omnibus):        0.039      Jarque-Bera (JB):
2.699
Skew:                 -0.070      Prob(JB):
0.259
Kurtosis:             2.111      Cond. No.          2.57
e+18
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 3.96e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

4.2 Removing insignificant variables

After removing outliers, insignificant predictors are removed one by one. Since `new_column` and `new_column2` have the same p-value and are highly similar,

`new_column2` is removed first.

```
In [10]: model_reduced1 = smf.ols(
          "Rate ~ Total + new_column",
          data=clean_table
        ).fit()

model_reduced2 = smf.ols(
          "Rate ~ Total",
          data=clean_table
        ).fit()

comparison = pd.DataFrame({
    "Model": [
        "Clean model",
        "Without new_column2",
        "Only Total"
    ],
    "Predictors": [
        "Total + new_column + new_column2",
        "Total + new_column",
        "Total"
    ],
    "Adj. R-squared": [
        model_clean.rsquared_adj,
        model_reduced1.rsquared_adj,
        model_reduced2.rsquared_adj
    ],
    "AIC": [
        model_clean.aic,
        model_reduced1.aic,
        model_reduced2.aic
    ],
    "BIC": [
        model_clean.bic,
        model_reduced1.bic,
        model_reduced2.bic
    ]
})

comparison
```

```
Out[10]:
```

	Model	Predictors	Adj. R-squared	AIC	BIC
0	Clean model	Total + new_column + new_column2	0.091573	245.799081	252.945161
1	Without new_column2	Total + new_column	0.091573	245.799081	252.945161
2	Only Total	Total	0.060778	247.498344	252.262397

```
In [11]: print(model_reduced1.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable:          Rate    R-squared:
0.115
Model:                  OLS     Adj. R-squared:
0.092
Method:                 Least Squares    F-statistic:
4.982
Date:                   Wed, 20 May 2026    Prob (F-statistic):      0.0
0923
Time:                   14:45:13    Log-Likelihood:      -11
9.90
No. Observations:      80    AIC:      2
45.8
Df Residuals:          77    BIC:      2
52.9
Df Model:               2
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.3899	0.622	3.843	0.000	1.152	
Total	-0.0001	4.57e-05	-2.492	0.015	-0.000	-2.29e-05
new_column	0.0077	0.004	1.909	0.060	-0.000	

```

=====
=====
Omnibus:                6.504    Durbin-Watson:
1.802
Prob(Omnibus):          0.039    Jarque-Bera (JB):
2.699
Skew:                   -0.070    Prob(JB):
0.259
Kurtosis:               2.111    Cond. No.      2.88e+04
=====
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.88e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [12]: print(model_reduced2.summary())
```

```

=====
                        OLS Regression Results
=====
=====
Dep. Variable:          Rate    R-squared:
0.073
Model:                  OLS    Adj. R-squared:
0.061
Method:                 Least Squares    F-statistic:
6.112
Date:                   Wed, 20 May 2026    Prob (F-statistic):      0.
0156
Time:                   14:45:14    Log-Likelihood:      -12
1.75
No. Observations:      80    AIC:      2
47.5
Df Residuals:          78    BIC:      2
52.3
Df Model:               1
Covariance Type:       nonrobust
=====
=====
                        coef      std err      t      P>|t|      [0.025      0.
975]
-----
Intercept      3.4670      0.266     13.040     0.000      2.938
3.996
Total      -0.0001     4.65e-05     -2.472     0.016     -0.000     -2.24
e-05
=====
=====
Omnibus:          8.285    Durbin-Watson:
1.918
Prob(Omnibus):    0.016    Jarque-Bera (JB):
3.039
Skew:             -0.064    Prob(JB):
0.219
Kurtosis:         2.054    Cond. No.      1.21
e+04
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.21e+04. This might indicate that there are strong multicollinearity or other numerical problems.

The reduced models are compared using adjusted R-squared, AIC, BIC, and predictor p-values. The preferred model should keep significant predictors while avoiding unnecessary duplicate variables.

5 Non-linear models

```
In [13]: nonLinear_model = smf.ols('Rate ~ Total + new_column + I(Total**2) + I(new_column**2)')
print(nonLinear_model.summary())
```

OLS Regression Results

```

=====
====
Dep. Variable:          Rate    R-squared:
0.128
Model:                  OLS     Adj. R-squared:
0.082
Method:                 Least Squares    F-statistic:
2.762
Date:                   Wed, 20 May 2026    Prob (F-statistic):          0.
0336
Time:                   14:45:14    Log-Likelihood:          -11
9.27
No. Observations:      80    AIC:          2
48.5
Df Residuals:          75    BIC:          2
60.5
Df Model:              4
Covariance Type:      nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025
Intercept	4.3173	2.124	2.033	0.046	0.086
Total	-0.0003	0.000	-1.390	0.169	-0.001
new_column	-0.0156	0.027	-0.567	0.573	-0.070
I(Total ** 2)	1.387e-08	1.78e-08	0.781	0.437	-2.15e-08
I(new_column ** 2)	7.848e-05	9.1e-05	0.862	0.391	-0.000

```

=====
====
Omnibus:              5.306    Durbin-Watson:
1.777
Prob(Omnibus):        0.070    Jarque-Bera (JB):
2.396
Skew:                 0.040    Prob(JB):
0.302
Kurtosis:             2.156    Cond. No.          7.40
e+08
=====
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 7.4e+08. This might indicate that there are strong multicollinearity or other numerical problems.

```

In [14]: logarithmic_model = smf.ols('Rate ~ np.log(Total) + np.log(new_column)',
print(logarithmic_model.summary())

```

OLS Regression Results

```

=====
=====
Dep. Variable:          Rate    R-squared:
0.128
Model:                  OLS     Adj. R-squared:
0.106
Method:                 Least Squares    F-statistic:
5.662
Date:                   Wed, 20 May 2026    Prob (F-statistic):      0.0
0508
Time:                   14:45:14    Log-Likelihood:      -11
9.28
No. Observations:      80    AIC:      2
44.6
Df Residuals:          77    BIC:      2
51.7
Df Model:              2
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025
Intercept	1.4389	3.045	0.473	0.638	-4.624
np.log(Total)	-0.4114	0.146	-2.816	0.006	-0.702
np.log(new_column)	0.9891	0.561	1.762	0.082	-0.129

```

=====
=====
Omnibus:                6.249    Durbin-Watson:
1.840
Prob(Omnibus):          0.044    Jarque-Bera (JB):
2.596
Skew:                   -0.011    Prob(JB):
0.273
Kurtosis:               2.118    Cond. No.
245.
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

6. Stepwise model selection

Forward stepwise starts with an empty model and adds predictors if they improve AIC. Backward stepwise starts with the full model and removes predictors if this improves AIC.

```
In [15]: candidate_predictors = ["Total", "new_column", "new_column2"]
response = "Rate"
```

```

def forward_stepwise(data, response, predictors):
    selected = []
    remaining = predictors.copy()
    current_model = smf.ols(f"{response} ~ 1", data=data).fit()
    current_aic = current_model.aic
    final_model = current_model

    while remaining:
        results = []

        for predictor in remaining:
            formula = response + " ~ " + " + ".join(selected + [predictor])
            model = smf.ols(formula, data=data).fit()
            results.append((model.aic, predictor, model))

        results.sort(key=lambda item: item[0])
        best_aic, best_predictor, best_model = results[0]

        if best_aic < current_aic:
            selected.append(best_predictor)
            remaining.remove(best_predictor)
            current_aic = best_aic
            final_model = best_model
            print("Added:", best_predictor, "AIC:", round(best_aic, 3))
        else:
            break

    return final_model, selected

def backward_stepwise(data, response, predictors):
    selected = predictors.copy()
    current_formula = response + " ~ " + " + ".join(selected)
    current_model = smf.ols(current_formula, data=data).fit()
    current_aic = current_model.aic

    while len(selected) > 1:
        results = []

        for predictor in selected:
            trial_predictors = selected.copy()
            trial_predictors.remove(predictor)
            formula = response + " ~ " + " + ".join(trial_predictors)
            model = smf.ols(formula, data=data).fit()
            results.append((model.aic, predictor, model, trial_predictors))

        results.sort(key=lambda item: item[0])
        best_aic, removed_predictor, best_model, best_predictors = results[0]

        if best_aic < current_aic:
            selected = best_predictors
            current_aic = best_aic
            current_model = best_model
            print("Removed:", removed_predictor, "AIC:", round(best_aic, 3))
        else:
            break

    return current_model, selected

```

```

In [16]: forward_model, forward_predictors = forward_stepwise(
        clean_table,

```

```
        response,  
        candidate_predictors  
    )  
  
    backward_model, backward_predictors = backward_stepwise(  
        clean_table,  
        response,  
        candidate_predictors  
    )  
  
    print("Forward selected predictors:", forward_predictors)  
    print("Backward selected predictors:", backward_predictors)
```

Added: Total AIC: 247.498

Added: new_column AIC: 245.799

Forward selected predictors: ['Total', 'new_column']

Backward selected predictors: ['Total', 'new_column', 'new_column2']

```
In [17]: print(forward_model.summary())  
         print(backward_model.summary())
```


OLS Regression Results

```

=====
=====
Dep. Variable:          Rate    R-squared:
0.115
Model:                  OLS     Adj. R-squared:
0.092
Method:                 Least Squares    F-statistic:
4.982
Date:                   Wed, 20 May 2026    Prob (F-statistic):      0.0
0923
Time:                   14:45:14    Log-Likelihood:      -11
9.90
No. Observations:      80    AIC:      2
45.8
Df Residuals:          77    BIC:      2
52.9
Df Model:               2
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	0.
Intercept	2.3899	0.622	3.843	0.000	1.152	
Total	-0.0001	4.57e-05	-2.492	0.015	-0.000	-2.29
new_column	0.0077	0.004	1.909	0.060	-0.000	

```

=====
=====
Omnibus:                6.504    Durbin-Watson:
1.802
Prob(Omnibus):          0.039    Jarque-Bera (JB):
2.699
Skew:                   -0.070    Prob(JB):
0.259
Kurtosis:               2.111    Cond. No.      2.88
e+04
=====
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.88e+04. This might indicate that there are strong multicollinearity or other numerical problems.

OLS Regression Results

```

=====
=====
Dep. Variable:          Rate    R-squared:
0.115
Model:                  OLS     Adj. R-squared:
0.092
Method:                 Least Squares    F-statistic:
4.982
Date:                   Wed, 20 May 2026    Prob (F-statistic):      0.0

```

```

0923
Time:                  14:45:14   Log-Likelihood:          -11
9.90
No. Observations:      80   AIC:                  2
45.8
Df Residuals:          77   BIC:                  2
52.9
Df Model:              2
Covariance Type:      nonrobust

```

```

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
Intercept      2.3899      0.622      3.843      0.000      1.152
3.628
Total      -0.0001    4.57e-05     -2.492      0.015     -0.000   -2.2
9e-05
new_column      0.0039      0.002      1.909      0.060     -0.000
0.008
new_column2     0.0039      0.002      1.909      0.060     -0.000
0.008
=====
=====
Omnibus:              6.504   Durbin-Watson:
1.802
Prob(Omnibus):        0.039   Jarque-Bera (JB):
2.699
Skew:                -0.070   Prob(JB):
0.259
Kurtosis:             2.111   Cond. No.            2.57
e+18
=====
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 3.96e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```

In [18]: stepwise_comparison = pd.DataFrame({
    "Model": ["Forward stepwise", "Backward stepwise"],
    "Predictors": [
        " + ".join(forward_predictors) if forward_predictors else "Intercept",
        " + ".join(backward_predictors) if backward_predictors else "Intercept"
    ],
    "Adj. R-squared": [
        forward_model.rsquared_adj,
        backward_model.rsquared_adj
    ],
    "AIC": [
        forward_model.aic,
        backward_model.aic
    ],
    "BIC": [
        forward_model.bic,
        backward_model.bic
    ]
})

```

```
    ]
  })

stepwise_comparison
```

Out [18]:

	Model	Predictors	Adj. R-squared	AIC	BIC
0	Forward stepwise	Total + new_column	0.091573	245.799081	252.945161
1	Backward stepwise	Total + new_column + new_column2	0.091573	245.799081	252.945161

The forward and backward models are compared using adjusted R-squared, AIC, and BIC. The preferred stepwise model is the one with lower AIC/BIC and a simpler set of predictors.

7. Final model comparison

All observed models are compared using adjusted R-squared, AIC, and BIC. A better model should have higher adjusted R-squared and lower AIC/BIC, while also remaining reasonably simple.

```
In [19]: final_comparison = pd.DataFrame({
    "Model": [
        "Initial model",
        "Cleaned model",
        "Reduced linear model",
        "Total-only model",
        "Quadratic model",
        "Logarithmic model",
        "Forward stepwise",
        "Backward stepwise"
    ],
    "Predictors": [
        "Total + new_column + new_column2",
        "Total + new_column + new_column2",
        "Total + new_column",
        "Total",
        "Total + new_column + Total^2 + new_column^2",
        "ln(Total) + ln(new_column)",
        " + ".join(forward_predictors),
        " + ".join(backward_predictors)
    ],
    "Adj. R-squared": [
        model.rsquared_adj,
        model_clean.rsquared_adj,
        model_reduced1.rsquared_adj,
        model_reduced2.rsquared_adj,
        nonLinear_model.rsquared_adj,
        logarithmic_model.rsquared_adj,
        forward_model.rsquared_adj,
        backward_model.rsquared_adj
    ],
    "AIC": [
```

```

        model.aic,
        model_clean.aic,
        model_reduced1.aic,
        model_reduced2.aic,
        nonLinear_model.aic,
        logarithmic_model.aic,
        forward_model.aic,
        backward_model.aic
    ],
    "BIC": [
        model.bic,
        model_clean.bic,
        model_reduced1.bic,
        model_reduced2.bic,
        nonLinear_model.bic,
        logarithmic_model.bic,
        forward_model.bic,
        backward_model.bic
    ]
})

final_comparison.sort_values("AIC")

```

Out[19]:

	Model	Predictors	Adj. R-squared	AIC	BIC
5	Logarithmic model	$\ln(\text{Total}) + \ln(\text{new_column})$	0.105556	244.558090	251.704170
1	Cleaned model	Total + new_column + new_column2	0.091573	245.799081	252.945161
2	Reduced linear model	Total + new_column	0.091573	245.799081	252.945161
6	Forward stepwise	Total + new_column	0.091573	245.799081	252.945161
7	Backward stepwise	Total + new_column + new_column2	0.091573	245.799081	252.945161
3	Total-only model	Total	0.060778	247.498344	252.262397
4	Quadratic model	Total + new_column + Total ² + new_column ²	0.081891	248.541844	260.451977
0	Initial model	Total + new_column + new_column2	0.036482	335.614034	343.211833

According to the comparison, the logarithmic model is selected as the best model. It has the lowest AIC/BIC and the highest adjusted R-squared among the tested models. Therefore, the final model explains **Rate** using **$\ln(\text{Total})$** and **$\ln(\text{new_column})$** .

8. Optional diagnostic tests for the selected model

```
In [20]: final_model = logarithmic_model

bp_test = het_breuschpagan(final_model.resid, final_model.model.exog)
gq_test = het_goldfeldquandt(final_model.resid, final_model.model.exog)
dw_value = durbin_watson(final_model.resid)

diagnostic_tests = pd.DataFrame({
    "Test": [
        "Breusch-Pagan heteroscedasticity test",
        "Goldfeld-Quandt heteroscedasticity test",
        "Durbin-Watson autocorrelation statistic"
    ],
    "Statistic": [
        bp_test[0],
        gq_test[0],
        dw_value
    ],
    "p-value": [
        bp_test[1],
        gq_test[1],
        np.nan
    ]
})

diagnostic_tests
```

```
Out[20]:
```

	Test	Statistic	p-value
0	Breusch-Pagan heteroscedasticity test	0.327313	0.849034
1	Goldfeld-Quandt heteroscedasticity test	0.825962	0.718091
2	Durbin-Watson autocorrelation statistic	1.839812	NaN

```
In [21]: # optional_task_interpretation
bp_pvalue = bp_test[1]
gq_pvalue = gq_test[1]

print("Optional task interpretation:")

if bp_pvalue < 0.05:
    print("Breusch-Pagan test: p-value is below 0.05, so heteroscedastici
else:
    print("Breusch-Pagan test: p-value is above 0.05, so there is no stro

if gq_pvalue < 0.05:
    print("Goldfeld-Quandt test: p-value is below 0.05, so heteroscedasti
else:
    print("Goldfeld-Quandt test: p-value is above 0.05, so there is no st

if 1.5 <= dw_value <= 2.5:
    print("Durbin-Watson statistic is close to 2, so there is no strong e
else:
    print("Durbin-Watson statistic is far from 2, so autocorrelation may
```

Optional task interpretation:

Breusch-Pagan test: p-value is above 0.05, so there is no strong evidence of heteroscedasticity.

Goldfeld-Quandt test: p-value is above 0.05, so there is no strong evidence of heteroscedasticity.

Durbin-Watson statistic is close to 2, so there is no strong evidence of autocorrelation.

For the selected logarithmic model, both heteroscedasticity tests have p-values above 0.05, so there is no strong evidence of heteroscedasticity. The Durbin-Watson statistic is close to 2, so there is no strong evidence of autocorrelation. Therefore, the optional diagnostic tests do not indicate serious violations for the final model.

The optional tests are used to check additional regression assumptions. For the Breusch-Pagan and Goldfeld-Quandt tests, a p-value below 0.05 indicates possible heteroscedasticity. The Durbin-Watson statistic is used to check autocorrelation; a value close to 2 suggests no strong autocorrelation.

9. Conclusions

The initial multiple regression model had low explanatory power and showed insignificant predictors. Residual diagnostics indicated possible influential observations, so outliers were detected using standardized residuals and Cook's distance. Outliers were removed one by one, and the model was rebuilt after every removal.

After removing outliers, insignificant and duplicate predictors were examined.

`new_column2` did not improve the model because it duplicated the information in `new_column`. Linear, quadratic, logarithmic, forward stepwise, and backward stepwise models were then compared.

The logarithmic model was selected as the best final model because it had the highest adjusted R-squared and the lowest AIC/BIC among the observed models. The final model is: `Rate ~ ln(Total) + ln(new_column)`.